

# I. DAVID REIN

817-800-0794 | [idadavidrein@gmail.com](mailto:idadavidrein@gmail.com) | [idadavidrein.github.io](https://github.com/idadavidrein) | [linkedin.com/in/idadavidrein](https://www.linkedin.com/in/idadavidrein) | [github.com/idadavidrein](https://github.com/idadavidrein)

## WORK & RESEARCH EXPERIENCE

---

**Model Evaluations and Threat Research (METR):** ML Research Scientist August 2024 - Present

- **Leading the General Autonomy team**, developing the science of measuring the capabilities of AI systems.

**NYU:** Research Scientist (Supervised by Sam Bowman) September 2022 - Present

- Research focused on scalable oversight methods and evaluations
- ***GPQA: A Graduate-Level Google-Proof Q&A Benchmark***: A new benchmark that has quickly become a standard evaluation for new foundation models (e.g. [Claude 3](#), [Llama 3](#), [Gemini 1.5 Pro](#), [Reka Flash](#))
- ***Debate Helps Supervise Unreliable Experts***: First positive results for [AI Safety via Debate](#) in a realistic setting.

**Cohere:** Member of Technical Staff September 2020 - November 2023

- Early employee, **co-founded the Embeddings team**, scaled up contrastive learning methods
- Trained large production LLMs on TPU v4 pods with Jax, designed and implemented core infrastructure.

**Duke University** January - August 2020

- ***Classification with Strategically Withheld Data***: AAAI 2021
- Evaluated the performance of a novel classification algorithm that is provably robust to strategically withheld data.
- Designed the graduate-level *Foundations of Deep Learning* (STA 790) course with Dr. David Banks

## EDUCATION

---

**Duke University** December 2021

- B.S. Computer Science, B.A. Philosophy - GPA: 3.81/4.0
- *Machine Learning (graduate level), Bayesian and Modern Statistics (graduate level), Applied Ethics, Metaphysics, Ethics and AI, Computational Microeconomics, Algorithms, Operating Systems, Computer Architecture, Science and Social Justice*

## LEADERSHIP

---

**Duke Undergraduate Machine Learning: Co-President** August 2018 - 2020

- Lead organizer of the 2019 Duke Datathon with 350+ participants; raised \$20k+ in sponsorship for the event.
- Hosted ~20 speakers from leading industry and research labs for seminars and workshops; average 20-40 attendees.

**Duke Effective Altruism (EA): Co-President** August 2019 - 2021

- Led and helped design the Arete Fellowship, a 12-week discussion-based program to introduce 20+ undergrads to EA.

## ACTIVITIES AND SKILLS

---

**Reinforcement Learning Implementations** May - June 2019

- Implemented REINFORCE (VPG), A2C with Generalized Advantage Estimation, and Proximal Policy Optimization.

**ASA Duke DataFest: 1st Place - Best Insight Award** April 2019

- Competed in a group against 425+ undergraduate (2/3) and graduate (1/3) students from 8 universities.
- Predicted fatigue from biometric data of the Canada women's rugby 7s team with a Cox proportional hazards model.

**Kenan Institute for Ethics Policy Prize in the Ethics of Emerging Tech: 2nd Place** April 2019

- Co-authored research paper on mechanics, ethics, and international policy of orbital debris and anti-satellite weaponry
- Presented paper at the 2019 Duke Conference on the Ethics of Emerging Technology.

**Languages and Tools:** Python, NumPy, TensorFlow, Jax, PyTorch, Scikit-Learn, Pandas